

p8106_hw1_lr3257

Leyang Rui

2025-02-15

p8106 Homework 1

Data import & cleaning

For easier model interpretation, I changed variables year_built and year_sold from categorical to numerical variables.

```
housing_test = read_csv("data/housing_test.csv") |>
  janitor::clean_names() |>
  mutate(year_built = as.numeric(year_built),
         year_sold = as.numeric(year_sold))

housing_train = read_csv("data/housing_training.csv") |>
  janitor::clean_names() |>
  mutate(year_built = as.numeric(year_built),
         year_sold = as.numeric(year_sold))
```


- (a) Fit a lasso model on the training data. Report the selected tuning parameter and the test error. When the 1SE rule is applied, how many predictors are included in the model?

```
x = model.matrix(sale_price ~ ., housing_train)[, -1]
y = housing_train[["sale_price"]]

lasso.model = glmnet(x, y, alpha = 1,
                     lambda = exp(seq(8, -5, length = 100)))

cv.lasso = cv.glmnet(x, y, alpha = 1,
                     lambda = exp(seq(8, -5, length = 100)))

plot(cv.lasso)
```



```
cv_lambda_min = cv.lasso$lambda.min

x_test = model.matrix(sale_price ~ ., housing_test)[, -1]
y_test = housing_test[["sale_price"]]

y_pred_lasso = predict(lasso.model, newx = x_test, s = cv_lambda_min)

test_mse_lasso = mean((y_test - y_pred_lasso)^2)
```

For a smallest test MSE of 4.407688×10^8 , the lambda is about 58.0094577

```
cv.lasso$lambda.1se
```

```
## [1] 474.1938
```

```
coef_1se = coef(cv.lasso, s = cv.lasso$lambda.1se)
num_coef_1se = sum(coef_1se[-1] != 0)
```

When the 1SE rule is applied, around 36 predictors are included in the model.

- (b) Fit an elastic net model on the training data. Report the selected tuning parameters and the test error. Is it possible to apply the 1SE rule to select the tuning parameters for elastic net? If the 1SE rule is applicable, implement it to select the tuning parameters. If not, explain why.

```

ctrl1 = trainControl(method = "cv", number = 10)

enet.model = train(sale_price ~ .,
  data = housing_train,
  method = "glmnet",
  tuneGrid = expand.grid(alpha = seq(0, 1, length = 21),
    lambda = exp(seq(-10, 10, length = 100))),
  trControl = ctrl1)
enet.model$bestTune

```

```

##      alpha  lambda
## 181  0.05 474.1938

```

```

x_test = housing_test
y_test = housing_test$sale_price

y_pred_enet = predict(enet.model, newdata = x_test)

test_mse_enet = mean((y_test - y_pred_enet)^2)

```

For a smallest test MSE of 4.3987576×10^8 , the lambda is about 474.1938297

Yes, the 1SE rule is applicable:

```

min_rmse = min(enet.model$results$RMSE)

se_rmse = sd(enet.model$results$RMSE) / sqrt(nrow(enet.model$resample))
lambda_1se_threshold = min_rmse + se_rmse

enet_1se = enet.model$results |>
  filter(RMSE <= lambda_1se_threshold) |>
  arrange(desc(lambda)) |>
  slice(1)

```

The tuning parameter lambda using 1SE rule is 9817.4745282

- (c) Fit a partial least squares model on the training data and report the test error. How many components are included in your model?

```

pls.model = pls(sale_price ~ ., data = housing_train,
  scale = TRUE, validation = "CV")
summary(pls.model)

```

```

## Data:      X dimension: 1440 39
## Y dimension: 1440 1
## Fit method: kernelpls
## Number of components considered: 39
##
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
##      (Intercept)  1 comps  2 comps  3 comps  4 comps  5 comps  6 comps
## CV           73685   33410   27986   25121   23916   23276   23140

```

```
## adjCV      73685    33405    27949    25049    23857    23217    23084
##           7 comps  8 comps  9 comps 10 comps 11 comps 12 comps 13 comps
## CV        23027    22997    23013    23009    22985    22974    22979
## adjCV      22974    22945    22957    22952    22928    22917    22922
##           14 comps 15 comps 16 comps 17 comps 18 comps 19 comps 20 comps
## CV        22978    22982    22976    22981    22992    22993    22996
## adjCV      22921    22925    22919    22924    22934    22935    22938
##           21 comps 22 comps 23 comps 24 comps 25 comps 26 comps 27 comps
## CV        22999    23001    23003    23003    23004    23004    23006
## adjCV      22940    22942    22944    22944    22945    22945    22947
##           28 comps 29 comps 30 comps 31 comps 32 comps 33 comps 34 comps
## CV        23007    23008    23008    23008    23008    23008    23008
## adjCV      22948    22948    22948    22948    22948    22948    22948
##           35 comps 36 comps 37 comps 38 comps 39 comps
## CV        23008    23008    23008    23008    23019
## adjCV      22948    22948    22948    22948    22933
##
## TRAINING: % variance explained
##           1 comps  2 comps  3 comps  4 comps  5 comps  6 comps  7 comps
## X           20.02    25.93    29.67    33.59    37.01    40.03    42.49
## sale_price   79.73    86.35    89.36    90.37    90.87    90.99    91.06
##           8 comps  9 comps 10 comps 11 comps 12 comps 13 comps 14 comps
## X           45.53    47.97    50.15    52.01    53.69    55.35    56.86
## sale_price   91.08    91.10    91.13    91.15    91.15    91.16    91.16
##           15 comps 16 comps 17 comps 18 comps 19 comps 20 comps
## X           58.64    60.01    62.18    63.87    65.26    67.10
## sale_price   91.16    91.16    91.16    91.16    91.16    91.16
##           21 comps 22 comps 23 comps 24 comps 25 comps 26 comps
## X           68.44    70.12    71.72    73.35    75.20    77.27
## sale_price   91.16    91.16    91.16    91.16    91.16    91.16
##           27 comps 28 comps 29 comps 30 comps 31 comps 32 comps
## X           78.97    80.10    81.83    83.55    84.39    86.34
## sale_price   91.16    91.16    91.16    91.16    91.16    91.16
##           33 comps 34 comps 35 comps 36 comps 37 comps 38 comps
## X           88.63    90.79    92.79    95.45    97.49    100.00
## sale_price   91.16    91.16    91.16    91.16    91.16    91.16
##           39 comps
## X           100.67
## sale_price   91.16
```

```
cv_mse = RMSEP(pls.model)
ncomp_cv = which.min(cv_mse$val[1,]) - 1

y_pred_pls = predict(pls.model, newdata = housing_test,
                     ncomp = ncomp_cv)
test_mse_pls = mean((y_test - y_pred_pls)^2)
```

The test MSE of the partial least squares model is 4.4962272×10^8 , while the model has 12 components

(d) Choose the best model for predicting the response and explain your choice.

```

lasso.model = train(sale_price ~ ., data = housing_train,
                    method = "glmnet",
                    tuneGrid = expand.grid(alpha = 1,
                                           lambda = exp(seq(-10, 10, length = 100))),
                    trControl = ctrl1)

pls.model = train(sale_price ~ ., data = housing_train,
                  method = "pls",
                  tuneLength = 10,
                  trControl = ctrl1)

resamp = resamples(list(lasso = lasso.model,
                        elastic_net = enet.model,
                        pls = pls.model))

summary(resamp)

```

```


##
## Call:
## summary.resamples(object = resamp)
##
## Models: lasso, elastic_net, pls
## Number of resamples: 10
##
## MAE
##           Min.   1st Qu.   Median     Mean   3rd Qu.     Max. NA's
## lasso      15459.28 16041.42 16478.26 16794.05 17178.92 18984.35    0
## elastic_net 15227.38 15931.63 16462.11 16608.21 17430.55 18344.76    0
## pls        20346.63 22287.51 22621.64 23344.97 25150.73 27392.32    0
##
## RMSE
##           Min.   1st Qu.   Median     Mean   3rd Qu.     Max. NA's
## lasso      20673.19 21336.58 21795.02 22930.58 22813.03 27968.86    0
## elastic_net 20700.39 21464.36 22754.18 22886.18 24283.02 25224.07    0
## pls        27985.84 29938.37 31771.50 32693.58 36357.34 38371.00    0
##
## Rsquared
##           Min.   1st Qu.   Median     Mean   3rd Qu.     Max. NA's
## lasso      0.8872972 0.8945839 0.8995071 0.9023868 0.9062259 0.9244404    0
## elastic_net 0.8672210 0.8895533 0.9077499 0.9013775 0.9162796 0.9233194    0
## pls        0.7339267 0.7872989 0.8063534 0.8031579 0.8241052 0.8543557    0

```

```

bwplot(resamp, metric = "RMSE")

```



I will choose the lasso model as the best model. Since the lasso model has both a smallest median RMSE and its distribution of RMSE is the most concentrated (despite the outliers), although it has a few outlier cases, it will still be the best and the most stable model in general.

- (e) If R package “caret” was used for the lasso in (a), retrain this model using R package “glmnet”, and vice versa. Compare the selected tuning parameters between the two software approaches. Should there be discrepancies in the chosen parameters, discuss potential reasons for these differences.

```
### Using caret here:
lasso.model = train(sale_price ~ ., data = housing_train,
                    method = "glmnet",
                    tuneGrid = expand.grid(alpha = 1,
                                           lambda = exp(seq(10, -10, length = 100))),
                    trControl = ctrl1)
lasso.model$bestTune
```

```
##      alpha  lambda
## 72      1 76.97143
```

```
cv_lambda_min
```

```
## [1] 58.00946
```

The results above show that the lambda from using the “caret” package is larger than that from using the “glmnet” package. The difference may come from how the “glmnet” method does cross validation together with building the model, while the “caret” method has an extra step of standardizing tuning process.